

```
In [48]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn import preprocessing
from IPython.display import display, HTML

import warnings
warnings.filterwarnings("ignore")
```

```
In [49]: df=pd.read_csv("D:/Final Project/archive (4)/2022/heart_data.csv")
display(HTML(df.head(10).to_html()))
```

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking
0	No	16.60	Yes	No	No	3.0	30.0	No
1	No	20.34	No	No	Yes	0.0	0.0	No
2	No	26.58	Yes	No	No	20.0	30.0	No
3	No	24.21	No	No	No	0.0	0.0	No
4	No	23.71	No	No	No	28.0	0.0	Yes
5	Yes	28.87	Yes	No	No	6.0	0.0	Yes
6	No	21.63	No	No	No	15.0	0.0	No
7	No	31.64	Yes	No	No	5.0	0.0	Yes
8	No	26.45	No	No	No	0.0	0.0	No
9	No	40.69	No	No	No	0.0	0.0	Yes

```
In [50]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 319795 entries, 0 to 319794
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   HeartDisease    319795 non-null   object 
 1   BMI              319795 non-null   float64
 2   Smoking          319795 non-null   object 
 3   AlcoholDrinking 319795 non-null   object 
 4   Stroke           319795 non-null   object 
 5   PhysicalHealth   319795 non-null   float64
 6   MentalHealth     319795 non-null   float64
 7   DiffWalking      319795 non-null   object 
 8   Sex               319795 non-null   object 
 9   AgeCategory      319795 non-null   object 
 10  Race              319795 non-null   object 
 11  Diabetic         319795 non-null   object 
 12  PhysicalActivity 319795 non-null   object 
 13  GenHealth        319795 non-null   object 
 14  SleepTime        319795 non-null   float64
 15  Asthma            319795 non-null   object 
 16  KidneyDisease    319795 non-null   object 
 17  SkinCancer        319795 non-null   object 
dtypes: float64(4), object(14)
memory usage: 43.9+ MB
```

In [51]: `duplicates = df[df.duplicated()]  
duplicates`

Out[51]:

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffW
<b>2182</b>	No	19.85	No	No	No	0.0	0.0	
<b>3182</b>	No	28.19	No	No	No	0.0	0.0	
<b>3397</b>	No	26.54	No	No	No	0.0	0.0	
<b>3650</b>	No	32.89	Yes	No	No	2.0	1.0	
<b>4061</b>	No	25.84	No	No	No	0.0	0.0	
...	...	...	...	...	...	...	...	...
<b>319671</b>	No	25.06	No	No	No	0.0	0.0	
<b>319689</b>	No	27.44	No	No	No	0.0	0.0	
<b>319726</b>	No	30.41	No	No	No	0.0	0.0	
<b>319751</b>	No	34.96	No	No	No	0.0	0.0	
<b>319776</b>	No	27.98	No	No	No	0.0	0.0	

18078 rows × 18 columns

In [52]: `df.drop_duplicates(inplace=True)  
df`

Out[52]:

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWk
0	No	16.60	Yes	No	No	3.0	30.0	
1	No	20.34	No	No	Yes	0.0	0.0	
2	No	26.58	Yes	No	No	20.0	30.0	
3	No	24.21	No	No	No	0.0	0.0	
4	No	23.71	No	No	No	28.0	0.0	
...	...	...	...	...	...	...	...	...
319790	Yes	27.41	Yes	No	No	7.0	0.0	
319791	No	29.84	Yes	No	No	0.0	0.0	
319792	No	24.24	No	No	No	0.0	0.0	
319793	No	32.81	No	No	No	0.0	0.0	
319794	No	46.56	No	No	No	0.0	0.0	

301717 rows × 18 columns

In [53]:

```
# Our Observation from the dataset
# Total Samle: 319795
# Total Duplicates : 18078
# The majority of features are categorical.
```

In [56]:

```
# Get unique values for the columns

unique_race = df['Race'].unique()
unique_general_health = df['GenHealth'].unique()
unique_diabetic = df['Diabetic'].unique()

print("Unique values in 'race' column:", unique_race)
print("Unique values in 'general health' column:", unique_general_health)
print("Unique values in 'diabetic' column:", unique_diabetic)
```

Unique values in 'race' column: ['3' '2' '1' '0' '5' '4']  
 Unique values in 'general health' column: ['3' '1' '2' '0' '4']  
 Unique values in 'diabetic' column: ['1' '0' '3' '4']

In [57]:

```
def categorize_bmi(BMI):
    if BMI <= 18.5:
        return 0
    elif 18.5 < BMI <= 25:
        return 1
    elif 25 < BMI <= 30:
        return 2
    elif BMI > 30:
        return 3

df['BMI'] = df['BMI'].apply(categorize_bmi)
```

In [58]:

```
df['AgeCategory']=df['AgeCategory'].replace('18-24','0')
df['AgeCategory']=df['AgeCategory'].replace('25-29','1')
```

```
df['AgeCategory']=df['AgeCategory'].replace('30-34','2')
df['AgeCategory']=df['AgeCategory'].replace('35-39','3')
df['AgeCategory']=df['AgeCategory'].replace('40-44','4')
df['AgeCategory']=df['AgeCategory'].replace('45-49','5')
df['AgeCategory']=df['AgeCategory'].replace('50-54','6')
df['AgeCategory']=df['AgeCategory'].replace('55-59','7')
df['AgeCategory']=df['AgeCategory'].replace('60-64','8')
df['AgeCategory']=df['AgeCategory'].replace('65-69','9')
df['AgeCategory']=df['AgeCategory'].replace('70-74','10')
df['AgeCategory']=df['AgeCategory'].replace('75-79','11')
df['AgeCategory']=df['AgeCategory'].replace('80 or older','12')
```

In [59]:

```
df['GenHealth']=df['GenHealth'].replace('Poor','0')
df['GenHealth']=df['GenHealth'].replace('Fair','1')
df['GenHealth']=df['GenHealth'].replace('Good','2')
df['GenHealth']=df['GenHealth'].replace('Very good','3')
df['GenHealth']=df['GenHealth'].replace('Excellent','4')

df['Race']=df['Race'].replace('American Indian/Alaskan Native','0')
df['Race']=df['Race'].replace('Asian','1')
df['Race']=df['Race'].replace('Black','2')
df['Race']=df['Race'].replace('White','3')
df['Race']=df['Race'].replace('Hispanic','4')
df['Race']=df['Race'].replace('Other','5')

df=df.replace('No','0')
df= df.replace('Yes','1')
df["Diabetic"]=df['Diabetic'].replace('No, borderline diabetes','3')
df["Diabetic"]=df['Diabetic'].replace('Yes (during pregnancy)','4')

df=df.replace('Female','0')
df= df.replace('Male','1')

display(HTML(df.head(10).to_html())))
```

	HeartDisease	BMI	Smoking	AlcoholDrinking	Stroke	PhysicalHealth	MentalHealth	DiffWalking
0	0	0	1	0	0	3.0	30.0	0
1	0	1	0	0	1	0.0	0.0	0
2	0	2	1	0	0	20.0	30.0	0
3	0	1	0	0	0	0.0	0.0	0
4	0	1	0	0	0	28.0	0.0	1
5	1	2	1	0	0	6.0	0.0	1
6	0	1	0	0	0	15.0	0.0	0
7	0	3	1	0	0	5.0	0.0	1
8	0	2	0	0	0	0.0	0.0	0
9	0	3	0	0	0	0.0	0.0	1


In [60]:

```
import seaborn as sns
#get correlations of each features in dataset
```

```

corrmat = df.corr()
top_corr_features = corrmat.index
plt.figure(figsize=(20,20))
#plot heat map
g=sns.heatmap(df[top_corr_features].corr(), annot=True, cmap="RdYlGn")

```



```

In [61]: import seaborn as sns
import matplotlib.pyplot as plt

columns_to_plot = [
    'HeartDisease', 'Smoking', 'AlcoholDrinking', 'Stroke', 'DiffWalking', 'Sex', 'Diabetic',
    'PhysicalActivity', 'GenHealth', 'SleepTime', 'Asthma', 'KidneyDisease', 'SkinCancer']

sns.set_style('whitegrid')

plt.figure(figsize=(20, 24))
for i, column in enumerate(columns_to_plot, 1):
    # Find the order for the current column by sorting its unique values

```

```
order = sorted(df[column].unique())

# Create the subplot
plt.subplot(4, 3, i)

# Generate the countplot with ordered X-axis
sns.countplot(x=column, data=df, order=order, palette='RdBu_r')

# Set the title and labels
plt.title(f'Count Plot of {column}')
plt.ylabel('Count')

# Adjust subplots layout
plt.tight_layout()
plt.show()
```



In [64]:

```

import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd



df['AgeCategory'] = pd.to_numeric(df['AgeCategory'])

# List of features to plot against 'HeartDisease'
features = ['Smoking', 'AlcoholDrinking', 'Stroke', 'DiffWalking', 'Sex',
            'Diabetic', 'PhysicalActivity', 'GenHealth', 'Asthma', 'KidneyDisease',
            'SkinCancer', 'AgeCategory', 'BMI']


# Setting the style
sns.set(style="whitegrid")

```


```
# Loop through the list of features
for feature in features:
    order = df[feature].value_counts().index.sort_values()
    plt.figure(figsize=(5, 3))
    sns.countplot(x=feature, data=df, hue='HeartDisease', order=order)
    plt.title(f'Count Plot of {feature} by HeartDisease')
    plt.legend(title='Heart Disease', labels=['No', 'Yes'])
    plt.show()
```




## Count Plot of Stroke by HeartDisease




## Count Plot of DiffWalking by HeartDisease




## Count Plot of Sex by HeartDisease




Count Plot of Diabetic by HeartDisease




Count Plot of PhysicalActivity by HeartDisease




Count Plot of GenHealth by HeartDisease




## Count Plot of Asthma by HeartDisease




## Count Plot of KidneyDisease by HeartDisease




## Count Plot of SkinCancer by HeartDisease



## Count Plot of AgeCategory by HeartDisease



## Count Plot of BMI by HeartDisease



```
In [13]: # There is a severe class imbalance (heart disease vs healthy)
# Imbalanced attributes : HeartDisease, Alcohol drinking, Stroke, DiffWalking, Race, Di
```

```
In [14]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df1 = df.copy(deep = True)

col = list(df.columns)
categorical_features = []
numerical_features = []
for i in df1.columns:
    if len(df[i].unique()) > 6:
        numerical_features.append(i)
    else:
        df1[i]=le.fit_transform(df1[i])
        categorical_features.append(i)
```

```
In [15]: df_corr=df.drop("Race", axis=1)


df_corr['HeartDisease']=df_corr['HeartDisease'].astype(float)
df_corr['Smoking']=df_corr['Smoking'].astype(float)
df_corr['AlcoholDrinking']=df_corr['AlcoholDrinking'].astype(float)
df_corr['Stroke']=df_corr['Stroke'].astype(float)
df_corr['DiffWalking']=df_corr['DiffWalking'].astype(float)
df_corr['Sex']=df_corr['Sex'].astype(float)
df_corr['AgeCategory']=df_corr['AgeCategory'].astype(float)
df_corr['Diabetic']=df_corr['Diabetic'].astype(float)
df_corr['PhysicalActivity']=df_corr['PhysicalActivity'].astype(float)
df_corr['GenHealth']=df_corr['GenHealth'].astype(float)
df_corr['Asthma']=df_corr['Asthma'].astype(float)
df_corr['KidneyDisease']=df_corr['KidneyDisease'].astype(float)
df_corr['SkinCancer']=df_corr['SkinCancer'].astype(float)
```

```
In [16]: df1.info()

<class 'pandas.core.frame.DataFrame'>
Index: 301717 entries, 0 to 319794
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   HeartDisease    301717 non-null   int32  
 1   BMI              301717 non-null   int64  
 2   Smoking          301717 non-null   int32  
 3   AlcoholDrinking 301717 non-null   int32  
 4   Stroke           301717 non-null   int32  
 5   PhysicalHealth   301717 non-null   float64 
 6   MentalHealth     301717 non-null   float64 
 7   DiffWalking      301717 non-null   int32  
 8   Sex               301717 non-null   int32  
 9   AgeCategory      301717 non-null   object  
 10  Race              301717 non-null   int32  
 11  Diabetic         301717 non-null   int32  
 12  PhysicalActivity 301717 non-null   int32  
 13  GenHealth        301717 non-null   int32  
 14  SleepTime        301717 non-null   float64 
 15  Asthma            301717 non-null   int32  
 16  KidneyDisease    301717 non-null   int32  
 17  SkinCancer        301717 non-null   int32  
dtypes: float64(3), int32(13), int64(1), object(1)
memory usage: 28.8+ MB
```

```
In [22]: colors = ["Pink",'SkyBlue','Yellow','#E94B3C','#2D2926']

corr=df_corr.corrwith(df_corr['HeartDisease']).sort_values(ascending=False).to_frame()
corr.columns=['Correlations']
plt.subplots(figsize=(5,10))
sns.heatmap(corr, annot=True, cmap=colors, linewidths=0.4, linecolor='black')
plt.title('Correlation with wrt outcome')
plt.show()
```



```
In [24]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, r
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
```

```

import seaborn as sns

X = df1.drop('HeartDisease', axis=1)
y = df1['HeartDisease']

# Split the dataset into the training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

# Feature scaling for Logistic regression
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Initialize the Logistic Regression model
logreg = LogisticRegression()
# Train the model
logreg.fit(X_train_scaled, y_train)
# Predict on the test set
y_pred_logreg = logreg.predict(X_test_scaled)

# Initialize the Random Forest model
rf = RandomForestClassifier(random_state=0)
# Train the model
rf.fit(X_train, y_train)
# Predict on the test set
y_pred_rf = rf.predict(X_test)

pos_label = 1

# Calculate metrics for Logistic Regression
imbalanced_logreg_metrics = {
    "Accuracy": accuracy_score(y_test, y_pred_logreg),
    "Precision": precision_score(y_test, y_pred_logreg, pos_label=pos_label),
    "Recall": recall_score(y_test, y_pred_logreg, pos_label=pos_label),
    "F1-Score": f1_score(y_test, y_pred_logreg, pos_label=pos_label),
    "AUC": roc_auc_score(y_test, logreg.predict_proba(X_test_scaled)[:, 1])
}

# Calculate metrics for Random Forest
imbalanced_rf_metrics = {
    "Accuracy": accuracy_score(y_test, y_pred_rf),
    "Precision": precision_score(y_test, y_pred_rf, pos_label=pos_label),
    "Recall": recall_score(y_test, y_pred_rf, pos_label=pos_label),
    "F1-Score": f1_score(y_test, y_pred_rf, pos_label=pos_label),
    "AUC": roc_auc_score(y_test, rf.predict_proba(X_test)[:, 1])
}

cm_logreg = confusion_matrix(y_test, y_pred_logreg)
cm_rf = confusion_matrix(y_test, y_pred_rf)

```

In [25]:

```
display(imbalanced_logreg_metrics)
display(cm_logreg)
```

```
{'Accuracy': 0.9105793450881612,
'Precision': 0.5340236686390533,
'Recall': 0.10560187216615474,
'F1-Score': 0.17633410672853828,
'AUC': 0.829571216956767}
```

```
array([[67963,    630],
       [ 6115,    722]], dtype=int64)
```

In [26]:

```
display(imbalanced_rf_metrics)
display(cm_rf)
```

```
{'Accuracy': 0.9011533872464537,
 'Precision': 0.36857749469214435,
 'Recall': 0.12695626736872898,
 'F1-Score': 0.18885987815491734,
 'AUC': 0.7821740575700927}
array([[67106,   1487],
       [ 5969,    868]], dtype=int64)
```

In [27]:

```
balanced_logreg = LogisticRegression(class_weight='balanced', random_state=0)
balanced_logreg.fit(X_train_scaled, y_train)
```

```
#train the balanced Random Forest model
balanced_rf = RandomForestClassifier(class_weight='balanced', random_state=0)
balanced_rf.fit(X_train, y_train)
```

```
# Make predictions with both balanced models
y_pred_balanced_logreg = balanced_logreg.predict(X_test_scaled)
y_pred_balanced_rf = balanced_rf.predict(X_test)
```

```
# Calculate performance metrics for both balanced models
balanced_metrics_logreg = {
```

```
    "Accuracy": accuracy_score(y_test, y_pred_balanced_logreg),
    "Precision": precision_score(y_test, y_pred_balanced_logreg),
    "Recall": recall_score(y_test, y_pred_balanced_logreg),
    "F1-Score": f1_score(y_test, y_pred_balanced_logreg),
    "AUC": roc_auc_score(y_test, balanced_logreg.predict_proba(X_test_scaled)[:, 1])
}
```

```
balanced_metrics_rf = {
    "Accuracy": accuracy_score(y_test, y_pred_balanced_rf),
    "Precision": precision_score(y_test, y_pred_balanced_rf),
    "Recall": recall_score(y_test, y_pred_balanced_rf),
    "F1-Score": f1_score(y_test, y_pred_balanced_rf),
    "AUC": roc_auc_score(y_test, balanced_rf.predict_proba(X_test)[:, 1])
}
```

In [28]:

```
display(balanced_metrics_logreg)
```

```
{'Accuracy': 0.7401034071324407,
 'Precision': 0.22519264712213183,
 'Recall': 0.7651016527716835,
 'F1-Score': 0.34796780416417217,
 'AUC': 0.8299289261450331}
```

In [29]:

```
display(balanced_metrics_rf)
```

```
{'Accuracy': 0.8674797825798753,
 'Precision': 0.23023057216054654,
 'Recall': 0.19716249817171275,
 'F1-Score': 0.21241727072171449,
 'AUC': 0.7573243336511979}
```

In [30]:

```
from imblearn.over_sampling import SMOTE
```

```
# Apply SMOTE to balance the training data
smote = SMOTE(random_state=0)
X_train_smote, y_train_smote = smote.fit_resample(X_train_scaled, y_train)

# Initialize and train the Logistic Regression model with SMOTE
smote_logreg = LogisticRegression(random_state=0)
smote_logreg.fit(X_train_smote, y_train_smote)

# Initialize and train the Random Forest model with SMOTE
smote_rf = RandomForestClassifier(random_state=0)
smote_rf.fit(X_train_smote, y_train_smote)

# Make predictions with both balanced models
y_pred_smote_logreg = smote_logreg.predict(X_test_scaled)
y_pred_smote_rf = smote_rf.predict(X_test_scaled)

# Calculate performance metrics for both balanced models
smote_metrics_logreg = {
    "Accuracy": accuracy_score(y_test, y_pred_smote_logreg),
    "Precision": precision_score(y_test, y_pred_smote_logreg),
    "Recall": recall_score(y_test, y_pred_smote_logreg),
    "F1-Score": f1_score(y_test, y_pred_smote_logreg),
    "AUC": roc_auc_score(y_test, smote_logreg.predict_proba(X_test_scaled)[:, 1])
}

smote_metrics_rf = {
    "Accuracy": accuracy_score(y_test, y_pred_smote_rf),
    "Precision": precision_score(y_test, y_pred_smote_rf),
    "Recall": recall_score(y_test, y_pred_smote_rf),
    "F1-Score": f1_score(y_test, y_pred_smote_rf),
    "AUC": roc_auc_score(y_test, smote_rf.predict_proba(X_test_scaled)[:, 1])
}
```

In [31]: `display(smote_metrics_logreg)`

```
{'Accuracy': 0.7401299217817844,
'Precision': 0.22497522299306244,
'Recall': 0.763639022963288,
'F1-Score': 0.34755691652243376,
'AUC': 0.8287304431049297}
```

In [32]: `display(smote_metrics_rf)`

```
{'Accuracy': 0.8791329709664589,
'Precision': 0.290979097909791,
'Recall': 0.23211935059236508,
'F1-Score': 0.2582377349279961,
'AUC': 0.7786520640545156}
```

In [ ]:

In [73]: `from imblearn.over_sampling import KMeansSMOTE`


```
KSMOTE = KMeansSMOTE(cluster_balance_threshold=0.1)
X_KSMOTE, y_KSMOTE = KSMOTE.fit_resample(X, y)

plt.figure(figsize = (5,3))
ax=sns.countplot( x = y_KSMOTE )
for container in ax.containers:
```

```

    ax.bar_label(container, label_type='center', rotation=0, color='white')
plt.title("Distribution After K_Means SMOTE ", size=14)
plt.show()

```



```

In [34]: from sklearn.model_selection import train_test_split
X_train , X_test, y_train, y_test = train_test_split(X_KSMOTE, y_KSMOTE , test_size=0.2, random_state=42)

scaler= StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

logreg_s=LogisticRegression(solver="liblinear", penalty="l2", C=0.00001, max_iter=1000)
logreg_s.fit(X_train_scaled,y_train)
y_pred_log=logreg_s.predict(X_test_scaled)

y_pred_proba_log = logreg_s.predict_proba(X_test_scaled)

KSMOTE_log_metrics = {
    "Accuracy": accuracy_score(y_test, y_pred_log),
    "Precision": precision_score(y_test, y_pred_log),
    "Recall": recall_score(y_test, y_pred_log),
    "F1-Score": f1_score(y_test, y_pred_log),
    "AUC": roc_auc_score(y_test, y_pred_proba_log[:, 1])
}

display(KSMOTE_log_metrics)

```

{'Accuracy': 0.9036865385596339,  
 'Precision': 0.9053528620709816,  
 'Recall': 0.9018331658876804,  
 'F1-Score': 0.9035895864790541,  
 'AUC': 0.9569976186746174}

```

In [74]: df_logreg = pd.DataFrame([imbalanced_logreg_metrics], index=['LogReg_Imbalanced'])
df_rf = pd.DataFrame([imbalanced_rf_metrics], index=['RF_Imbalanced'])
df_balanced_logreg = pd.DataFrame([balanced_metrics_logreg], index=['LogReg_Balanced'])
df_imbalanced_rf = pd.DataFrame([balanced_metrics_logreg], index=['RF_Balanced'])
df_smote_logreg = pd.DataFrame([smote_metrics_logreg], index=['LogReg_Smote'])

```

```
df_smote_rf = pd.DataFrame([smote_metrics_rf], index=['RF_Smote'])
df_KSMOTE_logreg = pd.DataFrame([KSMOTE_log_metrics], index=['LogReg_KSmote'])

simple_model = pd.concat([df_logreg, df_rf, df_balanced_logreg, df_imbalanced_rf, df_smote_rf])
compare_model = simple_model.transpose()

compare_model
```

Out[74]:

	LogReg_Imbalanced	RF_Imbalanced	LogReg_Balanced	RF_Balanced	LogReg_Smote	RF_Smote
Accuracy	0.910579	0.901153	0.740103	0.740103	0.740130	0.8791
Precision	0.534024	0.368577	0.225193	0.225193	0.224975	0.2909
Recall	0.105602	0.126956	0.765102	0.765102	0.763639	0.2321
F1-Score	0.176334	0.188860	0.347968	0.347968	0.347557	0.2582
AUC	0.829571	0.782174	0.829929	0.829929	0.828730	0.7786

In [39]: `import pickle`In [40]: `projectfile = 'trained_model.sav'`  
`pickle.dump(logreg_s, open(projectfile, 'wb'))`In [69]: `loaded_model = pickle.load(open('trained_model.sav', 'rb'))`